

Antonio J Cabrera
Paul Gazel-Anthoine
PAR1401

Lab 1: Experimental setup and tools

Spring 2018-19

Index

Session 1:

| | |
|--|---|
| 1.1 Node architecture and memory | 3 |
| 1.2 Serial compilation and execution | 4 |
| 1.3 Compilation and execution of OpenMP programs | 4 |
| 1.4 Strong vs Weak scalability | 5 |

Session 2

| | |
|---|---|
| 2.3 Exploring new task decompositions for 3DFFT | 9 |
|---|---|

Session 3

| | |
|-------------------------------------|----|
| 3.2.1 Initial version | 13 |
| 3.2.2 Improving ϕ | 16 |
| 3.2.3 Reducing parallelism overhead | 18 |

| | |
|-------------|----|
| Conclusions | 20 |
|-------------|----|

1.1 – Node architecture and memory

| | boada-1 to boada-4 | boada-5 | boada-6 to boada-8 |
|------------------------------------|--------------------|---------|--------------------|
| Number of sockets per node | 1 | 1 | 1 |
| Number of cores per socket | 6 | 6 | 8 |
| Number of threads per core | 2 | 2 | 1 |
| Maximum core frequency | 2.4 GHz | 2.6 GHz | 1.7 GHz |
| L1-I cache size (per-core) | 32 KB | 32 KB | 32 KB |
| L1-D cache size (per-core) | 32 KB | 32 KB | 32 KB |
| L2 cache size (per-core) | 256 KB | 256 KB | 256 KB |
| Last-level cache size (per-socket) | 12 MB | 15 MB | 20 MB |
| Main memory size (per-socket) | 12 GB | 31 GB | 16 GB |
| Main memory size (per-node) | 12 GB | 31 GB | 16 GB |

Table 1: Boada Architecture

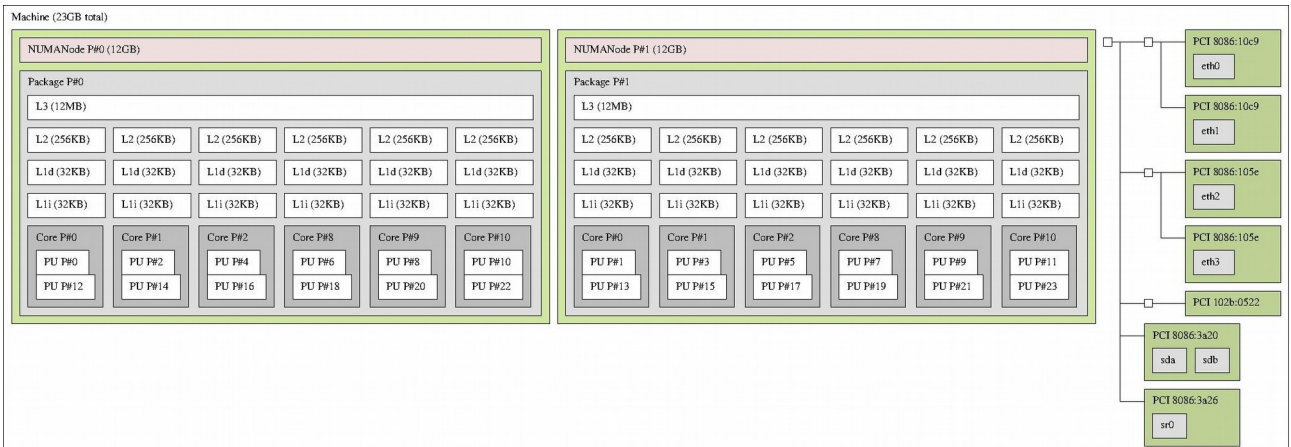


Figure 1: Boada-1 Architectural diagram

The architecture shown in figure 1 will allow us to work in parallel with 12 cores (24 threads) with two levels of cache each (64KB L1 + 256KB L2).

Each group of 6 cores share a cache memory (L3) of 12MB and a main memory of 12GB.

1.2 – Serial compilation and execution

In the `pi_seq.c` file, the function invocation and data structures used to measure execution time are `getusec()`, which uses the `timeval` struct to get the time of the day at a given time. By getting the time at the beginning of the program execution and at the end, it can compute the time spent in execution.

The binary is executed with two parameters (program and size) or else it will show the Usage guide. The execution is run through GNU `usr/bin/time`.

1.3 – Compilation and execution of OpenMP programs

The new lines indicate to the compiler which part of the code we want to execute in parallel. Each thread has a private variable `x` that gets its value by using the thread number. All `x`'s are used to compute the accumulated sum which are eventually summed via reduction as indicated in the 'pragma' instruction.

Results:

- 1 thread time elapsed → 3.96 seconds (3.94 user, 0.00 system) at 99% CPU
- 8 threads time elapsed → 0.63 seconds (4.84 user, 0.04 system) at 775% CPU

`export OMP_NUM_THREADS=$3` → It gathers the information at the third parameter and sets the environment variable `OMP_NUM_THREADS` with that value. That variable is later read inside the program.

| | 1 thread | 8 threads |
|-------------|----------------|----------------|
| Interactive | 3.9434 seconds | 0.5344 seconds |
| Queued | 3.9431 seconds | 0.5315 seconds |

Table 2: Pi computation, execution time comparison

As seen in Table 2, there was not a major difference between interactive and queued executions. We believe this data is so similar because at the time of our testing, `boada-1` was not heavily used by other students. We know that, in general, executing interactively while other people are using the same machine is not a good idea, since it would give skewed results because of some resources being used; testing is only to be done using the queued version.

1.4 – Strong vs Weak scalability

Strong scalability refers to how execution time is affected when adding more processors to deal with a constant amount of work.

Weak scalability refers to how execution time is affected when increasing proportionally the amount of work and the number of processors.

Strong Scalability (boada-4)




Figure 2: Time - Strong scalability (boada-4)

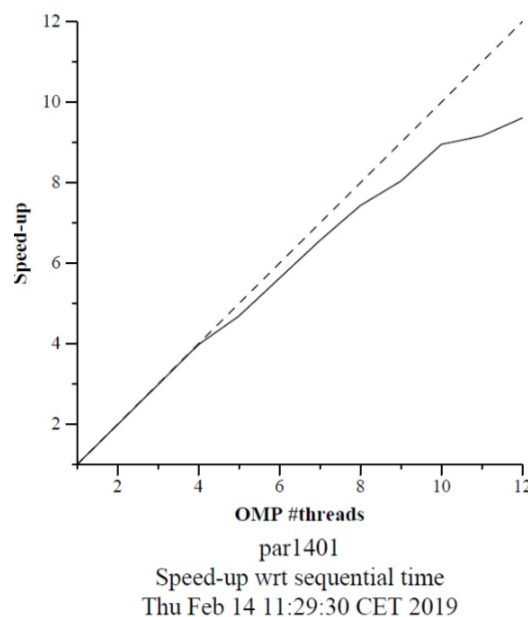


Figure 3: Speedup - Strong scalability (boada-4)

In figures 2 and 3, we can see that strong scalability for this program is almost ideal for up to 8 threads. After that, looking at figure 3, the line becomes flatter.

Weak Scalability (boada-3)




Figure 4: Speedup - Weak scalability (boada-4)

In figure 4, we can see that execution time is very stable, although not entirely flat. This means we lose some efficiency when we increase the size of the problem and the number of processors proportionally.

Strong Scalability (boada-5)




Figure 5: Time - Strong scalability (boada-5)

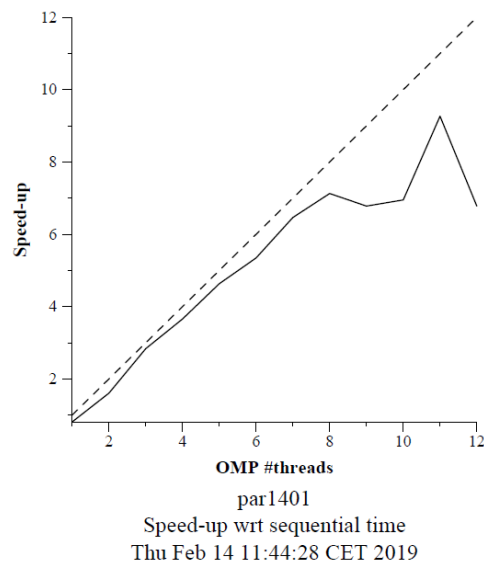


Figure 6: Speedup - Strong scalability (boada-5)

Similarly to boada-4, for up to 8 threads, execution time sees a quite ideal drop when increasing the number of processors. After that, when we get close to 0 seconds, the line flattens.

Weak Scalability (boada-5)




Figure 7: Speedup - Weak scalability (boada-5)

In figure 7, we see it is a bit of a mess and we cannot predict how execution time will evolve, although from 11 threads on, it seems that efficiency drops dramatically. It looks very efficient for a number of threads that follows this rule:

$$\text{threads} = 4k+2$$

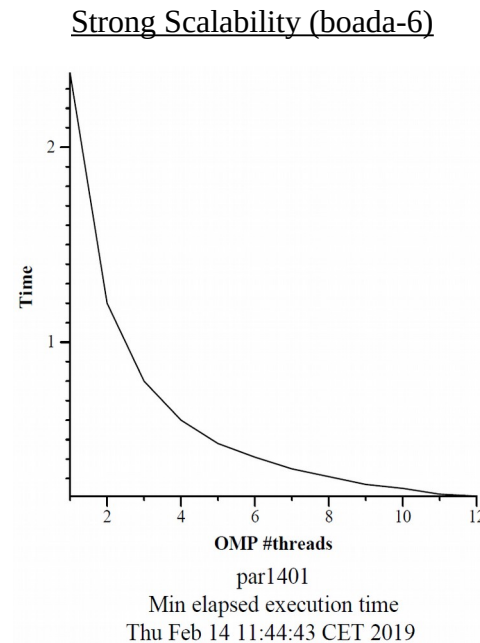


Figure 8: Time - Strong scalability (boada-6)




Figure 9: Speedup - Strong scalability (boada-6)

Figures 8 and 9 show that boada-6 follows the absolute ideal regarding strong scalability for this example.

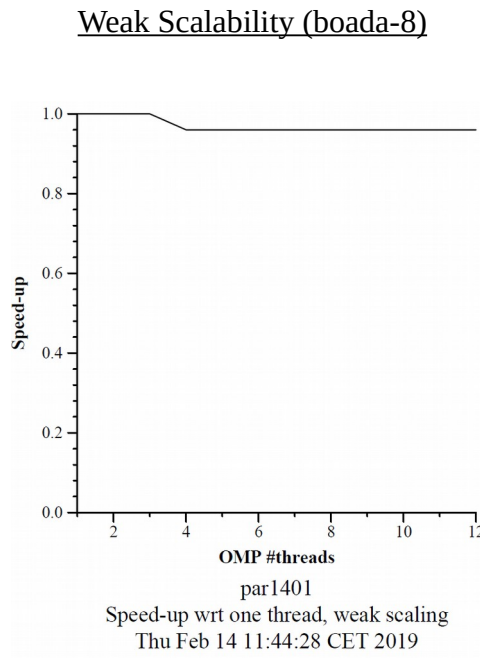


Figure 10: Speedup - Weak scalability (boada-8)

According to figure 10, boada-8 also behaves close to ideally regarding weak scalability.

Systematically analysing task decompositions with tareador

2.3 – Exploring new task decompositions for 3DFFT

| | T1 | T ∞ | Parallelism |
|------------|--------|------------|-------------|
| Sequential | 639780 | 639780 | 1 |
| V1 | 639780 | 639760 | 1,00003 |
| V2 | 639780 | 361525 | 1,77 |
| V3 | 639780 | 155065 | 4,13 |
| V4 | 639780 | 64750 | 9,88 |
| V5 | 639780 | 7861 | 81,39 |

Table 3: Parallelism study

By looking at table 3, we can already see a good level of parallelism for versions 4 and, specially, 5. This means that, if we have enough processors available, we will get a great increase in performance.

Parallel execution for different number of processors (V4)

Potential Strong Scalability (V4)




Figure 11: V4 time study

Potential Strong Scalability (V4)

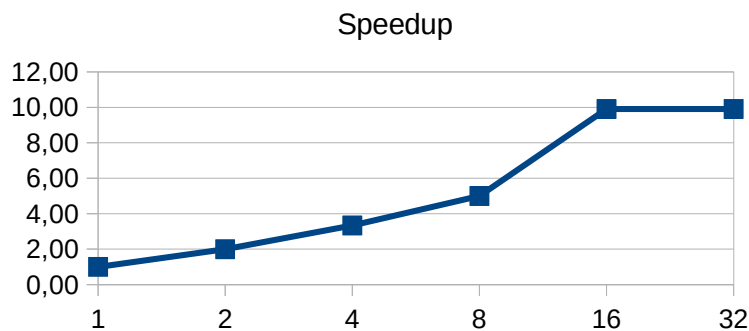


Figure 12: V4 speedup study

| V4 | | |
|------------|----------------|---------|
| Processors | Exec Time (ms) | Speedup |
| 1 | 639,8 | 1,00 |
| 2 | 320,3 | 2,00 |
| 4 | 191,9 | 3,33 |
| 8 | 128,0 | 5,00 |
| 16 | 64,6 | 9,90 |
| 32 | 64,6 | 9,90 |

Figure 13: V4 Times and Speedup

The limit is N, which gives the number of iterations and maximum number of threads than can be working in parallel if we decide to create a task for a loop.

Once we reach 10 threads, we reach said limit.

Parallel execution for different number of processors (V5)

Potential Strong Scalability (V5)

Finer granularity




Figure 14: V5 time study

Potential Strong Scalability (V5)

Finer granularity Speedup




Figure 15: V5 speedup study

| V5 | | |
|------------|----------------|---------|
| Processors | Exec Time (ms) | Speedup |
| 1 | 639,8 | 1,00 |
| 2 | 320,3 | 2,00 |
| 4 | 161,5 | 3,96 |
| 8 | 83,1 | 7,70 |
| 16 | 44,6 | 14,35 |
| 32 | 25,6 | 24,99 |

Figure 16: V5 Times and Speedup

Compared to V4, it is worth it to go to this granularity level if our machine is capable of working with 4 processors. In any case, the speedup increase is more obvious when we have 32 processors or more.

In this case, $N = 100$, so the speedup will not improve after we reach that number of processors.

We can scale to a higher number of processors compared to V4 thanks to the following code:

For each function with 3 nested loops:

| | |
|--------------------------|-------------------|
| for (k = 0; k < N; k++) | v4 version |
| tareador_start_task("A") | v5 version |
| for (j = 0; j < N; j++) | |
| tareador_start_task("B") | |
| for (i = 0; i < N; i++) | |
| // some computation | |
| tareador_end_task("B") | |
| tareador_end_task("A") | |

V4 only has task "A", while V5 has only task "B". This allows V5 to subdivide even more the problem (higher granularity) and take advantage of having more processors to work with.




Figure 17: V4 - Dependence graph




Figure 18: V5 - Dependence graph




Figure 19: V5 - Dependence graph (zoom-in)

Comparing figures 17 and 18, we can see that they have the same amount of tasks but a great increase in granularity.

It is worth it to have this level of granularity when you have more than N (loop size) processors. If that happens, subdividing at the nested loop level increases the parallelism potential.

Understanding the execution of OpenMP programs

3. After analyzing the performance (see next sections), we can say the scalability is not appropriate.

3.2.1 – Initial Version

1)

We obtain T_{par} with the configuration ‘Parallel functions duration’.

$T_{par} = 1,5315s$

$$T_{seq} = T_1 - T_{par} = 2,54 - 1,53 = 1,0098s$$

$$\text{Parallel fraction } (\phi) = 0,6026$$



Figure 20: Paraver obtained applying 'Parallel functions duration' configuration (1 thread)

2)

$$T_8 = 1,4744s$$

$$S_8 = 1,7097s$$




Figure 21: Paraver obtained applying 'Parallel functions duration' configuration (8 threads)

| | Running | Not created | Synchronization | Scheduling And Fork/Join | I/O | Others |
|--------------|---------|-------------|-----------------|-----------------------------|------|--------|
| THREAD 1.1.1 | 95.94 | 0.00 | 3.94 | 0.10 | 0.02 | 0.00 |
| THREAD 1.1.2 | 37.69 | 61.32 | 0.98 | 0.00 | 0.02 | 0.00 |
| THREAD 1.1.3 | 37.38 | 61.32 | 1.29 | 0.00 | 0.02 | 0.00 |
| THREAD 1.1.4 | 34.40 | 61.32 | 4.27 | 0.00 | 0.02 | 0.00 |
| THREAD 1.1.5 | 34.56 | 61.32 | 4.11 | 0.00 | 0.02 | 0.00 |
| THREAD 1.1.6 | 37.37 | 61.32 | 1.30 | 0.00 | 0.02 | 0.00 |
| THREAD 1.1.7 | 34.93 | 61.32 | 3.73 | 0.00 | 0.02 | 0.00 |
| THREAD 1.1.8 | 37.57 | 61.32 | 1.10 | 0.00 | 0.02 | 0.00 |
| Total | 349.83 | 429.23 | 20.71 | 0.10 | 0.13 | 0.00 |
| Average | 43.73 | 61.32 | 2.59 | 0.10 | 0.02 | 0.00 |
| Maximum | 95.94 | 61.32 | 4.27 | 0.10 | 0.02 | 0.00 |
| Minimum | 34.40 | 61.32 | 0.98 | 0.10 | 0.02 | 0.00 |
| Stddev | 19.78 | 0.00 | 1.43 | 0.00 | 0.00 | 0.00 |
| Avg/Max | 0.46 | 1.00 | 0.61 | 1.00 | 0.92 | 1.00 |

Figure 22: Percentage of time spent in the different OpenMP states (8 threads)

3)

$T1 = 2,1065s$

$T8 = 1,1410s$

$S8 = 1,846$

4)





Figure 23: Strong Scalability plots (Initial 3dfft_omp)

The ideal S_∞ is 2,51 (by Amhdal's law) and for S8 we get 1,85. However, watching at the plots in figure 23, we see it's difficult to significantly increase the speedup. Thus, we assume there is some overhead that will not be reduced by parallelism and estimate the actual upper limit somewhere around 2.

3.2.2 – Improving ϕ

The function causing the low value for ϕ is 'init_complex_grid'.

$T_1 = 2,3835s$

$T_8 = 1,0023s$

$T_{par} = 2,1363$

$\phi = 0,896$

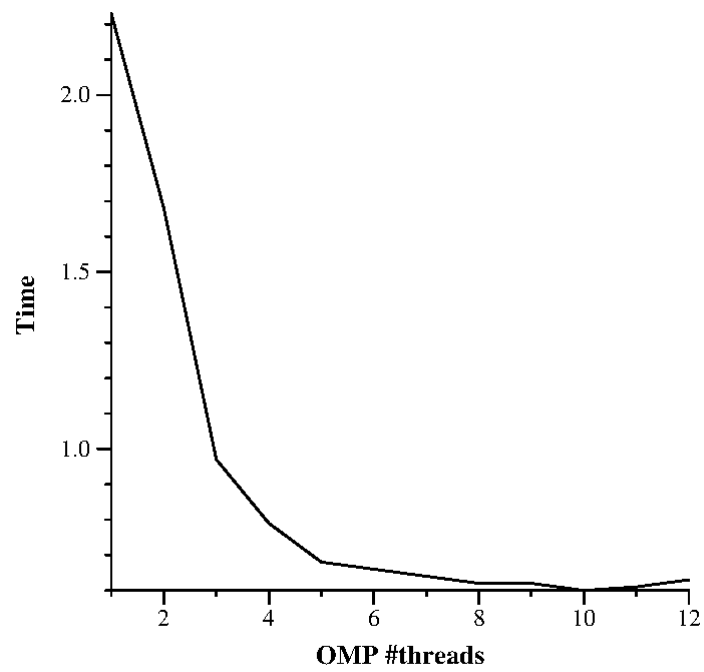
$S_8 = 2,378$



Figure 24: Paraver obtained applying 'Parallel functions duration' configuration (1 thread)

| | Running | Not created | Synchronization | Scheduling and Fork/Join | I/O | Others |
|----------------|---------|-------------|-----------------|--------------------------|------|--------|
| THREAD 1.1.1 | 96.95 | 0.00 | 2.90 | 0.11 | 0.03 | 0.00 |
| THREAD 1.1.2 | 56.83 | 39.98 | 3.17 | 0.00 | 0.03 | 0.00 |
| THREAD 1.1.3 | 57.96 | 39.97 | 2.03 | 0.00 | 0.03 | 0.00 |
| THREAD 1.1.4 | 57.36 | 39.98 | 2.63 | 0.00 | 0.03 | 0.00 |
| THREAD 1.1.5 | 57.78 | 39.98 | 2.21 | 0.00 | 0.03 | 0.00 |
| THREAD 1.1.6 | 57.12 | 39.99 | 2.87 | 0.00 | 0.03 | 0.00 |
| THREAD 1.1.7 | 57.99 | 39.98 | 2.00 | 0.00 | 0.03 | 0.00 |
| THREAD 1.1.8 | 56.88 | 40.00 | 3.09 | 0.00 | 0.03 | 0.00 |
| Total | 498.88 | 279.88 | 20.90 | 0.11 | 0.23 | 0.00 |
| Average | 62.36 | 39.98 | 2.61 | 0.11 | 0.03 | 0.00 |
| Maximum | 96.95 | 40.00 | 3.17 | 0.11 | 0.03 | 0.00 |
| Minimum | 56.83 | 39.97 | 2.00 | 0.11 | 0.03 | 0.00 |
| Stdev | 13.08 | 0.01 | 0.44 | 0.00 | 0.00 | 0.00 |
| Avg/Max | 0.64 | 1.00 | 0.82 | 1.00 | 0.91 | 1.00 |

Figure 25: Percentage of time spent in the different OpenMP states (8 threads)



par1401

Min elapsed execution time

Tue Mar 5 15:20:46 CET 2019




Figure 26: Strong Scalability plots (Improved ϕ)

Similarly to 3.2.1.4, if we compute the ideal speedup for infinite processors, we should achieve an approximated speedup of 9,6. However, looking at the plots in figure 26, we estimate the upper bound to be closer to 4 due to expected overheads.

3.2.3 – Reducing parallelism overheads

$T_1 = 2,5116s$

$T_8 = 0,8836s$

$T_{par} = 2,2467$

$\phi = 0,8945$

$S_8 = 2,8425$




Figure 27: Paraver obtained applying 'Parallel functions duration' configuration (1 thread)

| | Running | Not created | Synchronization | Scheduling and Fork/Join | I/O | Others |
|----------------|---------|-------------|-----------------|--------------------------|------|--------|
| THREAD 1.1.1 | 99.20 | 0.00 | 0.71 | 0.09 | 0.00 | 0.00 |
| THREAD 1.1.2 | 41.34 | 51.76 | 6.90 | 0.00 | 0.00 | 0.00 |
| THREAD 1.1.3 | 40.02 | 51.76 | 8.21 | 0.00 | 0.00 | 0.00 |
| THREAD 1.1.4 | 45.19 | 51.76 | 3.05 | 0.00 | 0.00 | 0.00 |
| THREAD 1.1.5 | 45.17 | 51.80 | 3.03 | 0.00 | 0.00 | 0.00 |
| THREAD 1.1.6 | 36.95 | 51.81 | 11.24 | 0.00 | 0.00 | 0.00 |
| THREAD 1.1.7 | 47.58 | 51.80 | 0.61 | 0.00 | 0.00 | 0.00 |
| THREAD 1.1.8 | 45.20 | 51.81 | 2.99 | 0.00 | 0.00 | 0.00 |
| Total | 400.64 | 362.52 | 36.74 | 0.09 | 0.01 | 0.00 |
| Average | 50.08 | 51.79 | 4.59 | 0.09 | 0.00 | 0.00 |
| Maximum | 99.20 | 51.81 | 11.24 | 0.09 | 0.00 | 0.00 |
| Minimum | 36.95 | 51.76 | 0.61 | 0.09 | 0.00 | 0.00 |
| Stdev | 18.84 | 0.02 | 3.55 | 0.00 | 0.00 | 0.00 |
| Avg/Max | 0.50 | 1.00 | 0.41 | 1.00 | 0.35 | 1.00 |

Figure 28: Percentage of time spent in the different OpenMP states (8 threads)

Finally, we obtain S_∞ of approximately 9,5 using Amhdal's law. In this case, by reducing overheads, we can get closer to that theoretical number.

The speedup plot in figure 29 doesn't allow for us to estimate an upper bound, but we know it is at least 6, which is a good 50% increase compared to the 3.2.2 results.





Figure 29: Strong Scalability plots (reduced overheads)

Summary: understanding the parallel execution of 3DFFT

| Version | ϕ | S_{∞} | T1 | T8 | S8 |
|--|--------|--------------|------|------|------|
| Initial version in 3dfft_omp.c | 0,6 | 2,51 | 2,54 | 1,47 | 1,71 |
| New version with improved ϕ | 0,9 | 9,61 | 2,38 | 1 | 2,38 |
| Final version with reduced parallelisation overheads | 0,89 | 9,48 | 2,51 | 0,88 | 2,84 |

*Table 4: Performance results of 3dfft***Conclusions**

In this sessions we have learned the basics of parallelism and some tools to evaluate potential parallelism.

We have been working with dependence graphs, Paraver, Strong and Weak scalability graphs.

Through a few examples, we have seen the importance of overheads in limiting the maximum parallelism attainable due to physical limitations, even though the theoretical calculation (Amhdal's law) aims higher.

It is really important that we understand these concepts so we can minimize overheads and reach the maximum potential in our parallel programs.